

 Review the assignment due date

Assignment 2 - SearchComplete

Assignment Objectives

1. Learn how to implement search algorithms in python
2. Learn how search algorithms can be used in practical application
3. Learning the differences between BFS, DFS, and UCS via implementation
4. Analyze the differences between search algorithms by comparing outputs
5. Learning how to build a search tree from textual data
6. Build a basic autocomplete feature that suggests words as the user types, using different search strategies.
7. Analyze how each algorithm affects the order and quality of suggestions, and learn when to choose each one.

Pre-Requisites

- **Basic Python:** Familiarity with Python syntax, data structures (lists, dictionaries, queues), and basic algorithms.
- **Search Algorithms:** Theoretical understanding of BFS, DFS, and UCS
- **Tree:** Prior knowledge of Tree data structures is helpful.
- **Data Structures:** High level understanding of Data Structures like Stacks, Queues, and Priority Queues is required.

Overview

Imagine you're an intern at a cutting-edge tech company called "WordWizard." Your first task: upgrade their revolutionary messaging app, "ChatCast," to include a mind-blowing autocomplete feature. The goal is simple – as users type, the app magically suggests the words they might be looking for, making conversations faster and more fun!

But here's the twist: Your quirky, genius boss, Dr. Lexico, insists on using classic search algorithms to power this futuristic feature. "Forget fancy neural networks," she exclaims. "Let's prove that good old BFS, DFS, and UCS can still deliver the goods!"

So, you're handed a massive dictionary of Gen Z slang and challenged to build the autocomplete engine. Can you master the algorithms, construct a word-filled tree, and unleash the power of search to create an autocomplete experience that will make even the most texting-savvy teen say, "OMG, this is lit!"?

The future of "ChatCast" (and your internship) depends on it. Time to dive into the code and become a word-suggesting wizard!

Lab Description

1. First step

- Clone the repo and run `main.py`

```
python main.py
```

- If you're on linux/mac and the former doesn't work for you

```
python3 main.py
```

2. Explore the Starter Code:

- Review the provided `Autocomplete` class. It handles building the tree from a text document, setting up a basic user interface, and providing a framework for the `suggest` method.

3. Implement Search Algorithms:

- Your main task is to complete the `suggest` methods. These methods should take a prefix as input and return a list of word suggestions.
- You'll implement multiple versions of `suggest`:
 - `suggest_bfs`: Breadth-First Search
 - `suggest_dfs`: Depth-First Search
 - `suggest_ucs`: Uniform-Cost Search

Background: Autocomplete as a Search Problem

Alright! Let's give you some context before you get into the weeds of the starter code. Autocomplete might seem like some complicated magic, but at its core, it's just an application of search algorithms on a tree (that's how it's done in this assignment for your simplicity, but it's done very differently in real word). Let's break down how this works:

The Search Space: A Tree of Characters

To implement the autocomplete feature, you would build a tree of characters, which will be the search space for this search problem. In your starter code, you're given a `document` (a `txt` file) of several words. Imagine each word in your document is broken down into its individual letters. Now, picture these letters arranged in a single tree-like structure, for example look at the tree diagram below:

Tree Diagram

For example, let the document that is given to you be -

```
air ball cat car card carpet carry cap cape
```

```
graph TD;
  ROOT-->A_air[A];
  A_air[A]-->I_air[I]
```

```

I_air[I]-->R_air[R]

ROOT-->B
B-->A_ball[A]
A_ball[A]-->L_ball1[L]
L_ball1[L]-->L_ball2[L]

ROOT-->C
C-->A_cat[A]
A_cat[A]-->T

A_cat[A]-->R
R-->D

R-->P_carpet[P]
P_carpet[P]-->E_carpet[E]
E_carpet[E]-->T_carpet[T]

R-->R_carry[R]
R_carry[R]-->Y

A_cat[A]-->P_ape[P]
P_ape[P]-->E_ape[E]

```

Above is a diagram of the tree that is build from the example **document** given above. Note how the *tree* starts with a common **root**

- This is what the search space for your search problem would look like.
- You will traverse the *tree* starting from the last node of the prefix that the user enters to generate autocomplete suggestions.

The Search Problem

When a user types a prefix (e.g., "ca"), the autocomplete feature needs to find all the words in the *tree* that start with that prefix. This translates to a search problem:

- **Initial state:** The node representing the last letter of the prefix ("a" in our example).
- **Action** - a transition between one letter to the next letter in the *tree*
- **Goal:** The end of the word(s) (that start with the given prefix) in the *tree*. Note how there could be multiple goals in this problem.
- **Path:** The sequence of characters from the root to a goal node represents a complete word.

Search Algorithms

We can employ various search algorithms to traverse this *tree* and find our goal nodes (complete words).

- **Breadth-First Search (BFS):** Explores the *tree* level-by-level, ensuring we find the shortest words first.
- **Depth-First Search (DFS):** Dives deep into the *tree*, potentially finding longer, less common words first.

- **Uniform-Cost Search (UCS):** Considers the frequency of each character transition to prioritize more likely words based on the prefix.

Multiple Goals and Paths

In autocomplete, we're not just looking for a single goal node. We want to find *all* the goal nodes (words) that follow from the prefix. Furthermore, we're interested in the entire path from the root to each goal node, as this path represents the complete suggested word.

Your Task:

Your task is to implement BFS, DFS, and UCS to traverse the *tree* and generate autocomplete suggestions. You'll see how different algorithms affect the order and type of words suggested, and understand the trade-offs involved in choosing one over the other.

Starter Code

For the starter code you have been given 3 files -

1. **autocomplete.py** - This is where all your code that you write will go.
2. **main.py** - This file is responsible to setting up and running the autocomplete feature. Modifying this file is optional. Feel free to use this file for debugging or playing around with the autocomplete feature.
3. **utilities.py** - This file contains the code to read the document provided and building the Graphical User Interface for the autocomplete feature. This file is not related to the core logic of the autocomplete feature. Please do not modify this file.

autocomplete.py

- This file has a **Node** class defined for you -
 - Each Node represents a single character within a word. The `Node` class has 1 attribute -
 1. **children** - This is a dictionary that stores -
 - Keys - Characters that which follow the current character in a word.
 - Values - **Node** objects, representing the next character in the sequence. **You might (most likely will) want the **Node** class keep track of more things depending on how you implement you **suggest** methods.**
- The file also has an **autocomplete** class defined for you -
 - The Engine Behind the Suggestions
 - **Attributes**
 - **root**: A root node of the tree. The tree stores all the words of the document in a tree structure, where each **Node** is character.
 - **Methods**
 - **__init__(document="")**:
 - Initializes an empty tree (the **root** node).
 - If a **document** string is provided, it builds the tree from that document.
 - document is a space separated textfile, example below.

- `air ball cat car card carpet carry cap cape`

- `build_tree(document)` #TODO:

- As the name of the function suggests, takes a text string `document` and builds a tree of words, where each `Node` is a character.
- The implementation of this method has been left up to you.

Student Tasks:

The main goal of the lab activity is for students to implement the `build_tree`, `suggest_bfs`, `suggest_ucs`, and `suggest_dfs` methods.

0. TODO: Intuition of the code written

- For all code that you will write for this assignment (which is not a lot), you must provide a brief intuition (1-2 sentences) of the major control structures of your code in the reports section at the bottom of this readme.
- You are not being asked to write a story, keep it concise and precise (remember, 1-2 sentences, at most 3).

Consider the `fizz-buzz` code given below:

```
def fizzbuzz(n):
    for i in range(1, n + 1):
        if i % 15 == 0:
            print("FizzBuzz")
        elif i % 3 == 0:
            print("Fizz")
        elif i % 5 == 0:
            print("Buzz")
        else:
            print(i)
```

Now this is what your explanation should (somewhat) look like -

Iterates through a range of numbers n printing that number unless the number is a multiple of 3 or 5 where instead "Fizz" or "Buzz" is printed respectively. "FizzBuzz" is printed if the number is a multiple of both 3 and 5.

1. TODO: `build_tree(document)`

[!NOTE] **TODO: Draw the tree diagram of test.txt given in the starter code** - Upload the image into your `readme` into the reports section in the end of this readme.

What it does:

- Takes a text `document` as input.

- Splits the document into individual words.
- Inserts each word into a tree (prefix tree) data structure.
- Each character of a word becomes a node in the tree.

Your task:

- Complete the `for` loop within the `build_tree` method.

2. TODO: `suggest_bfs(prefix)`**What it does:**

- Implements the Breadth-First Search (BFS) algorithm on the tree.
- Takes a `prefix` (the letters the user has typed so far) as input.
- Finds all words in the tree that start with the `prefix`.

Your task:

- Start from the node that corresponds to the last character of the `prefix`.
- Using BFS traverse the sub tree and build a list of suggestions.
- **Run your code with the `genZ.txt` file and `suggest_bfs()` method that you just implemented with the prefix `"th"` and note the the autocompleted suggestions it generates in the *Reports Section* below. Make sure you note down the suggestions in the same order in which they are originally displayed on your screen.**

3. TODO: `suggest_dfs(prefix)`**What it does:**

- Implements the Depth-First Search (DFS) algorithm on the tree.
- Takes a `prefix` as input.
- Finds all words in the tree that start with the `prefix`.

Your task:

- Start from the node that corresponds to the last character of the `prefix`.
- Using DFS traverse the sub tree and build a list of suggestions.
- **Explain your intuition in recursive DFS VS stack-based DFS, and which one you used. Write this in the section provided at the end of this readme.**
- **Run your code with the `genZ.txt` file and `suggest_dfs()` method that you just implemented with the prefix `"th"` and note the the autocompleted suggestions it generates in the *Reports Section* below. Make sure you note down the suggestions in the same order in which they are originally displayed on your screen.**

4. TODO: `suggest_ucs(prefix)`**What it does:**

- Implements the Uniform Cost Search (UCS) algorithm on the tree.
- Takes a `prefix` as input.
- Finds all words in the tree that start with the `prefix`.

- Prioritizes suggestions based on the frequency of characters appearing after previous characters.

Your task:

- Update `build_tree()` to store the path cost. The path cost is the inverse frequencies of that letter/char following that prefix of characters.
 - Using the inverse of these frequencies creates a lower path cost for more frequent character sequences.
- Start from the node that corresponds to the last character of the `prefix`.
- Using UCS traverse the sub tree and build a list of suggestions.
- **Run your code with the `genZ.txt` file and `suggest_ucs()` method that you just implemented with the prefix `"th"` and note the the autocompleted suggestions it generates in the *Reports Section* below. Make sure you note down the suggestions in the same order in which they are originally displayed on your screen.**

[!NOTE] This is not optional Try experimenting with different approaches and compare the results! Try typing different prefixes in the GUI and observe how the suggested words change depending on which search algorithm you're using. This will help you gain a deeper understanding of their strengths and weaknesses.

Note down these observations in the reports section provided at the end of this readme

What to Submit

1. **Completed `autocomplete.py` file:** Containing your implementations of the `build_tree`, `suggest_bfs`, `suggest_dfs`, and `suggest_ucs` methods.
2. **Completed *Reports Section* at the bottom of the `readme.md` file:** Briefly explaining wherever necessary, and completing the required tasks in the *Reports Section*.

Rubric

Criteria	Points (Example)
Diagram and explanation for <code>build_tree</code>	10%
Correctness of <code>build_tree</code>	10%
Explanation of <code>build_tree</code>	10%
Correctness of <code>suggest_bfs</code>	10%
Explanation of <code>suggest_bfs</code>	10%
Correctness of <code>suggest_dfs</code>	10%
Explanation of <code>suggest_dfs</code>	10%
Correctness of <code>suggest_ucs</code>	10%
Explanation of <code>suggest_ucs</code>	10%
Experimentation	10 %

A Reports section

383GPT

Did you use 383GPT at all for this assignment (yes/no)? Yes!

build_tree

Tree diagram

```
graph TD;
  ROOT-->T
  T-->H

  H-->E_there[E]-->R_there[R]-->E_there2[E]-->end_there[END]

  E_there[E]-->end_the[END]

  E_there[E]-->E_thee[E]-->end_thee[END]

  E_there[E]-->I_their[I]-->R_their[R]-->end_their[END]

  H-->A_that[A]-->T_that[T]-->end_that[END]

  A_that[A]-->G_thag[G]-->end_thag[END]

  H-->O_thou[O]-->U_though[U]-->G_though[G]-->H_though[H]-->end_though[END]

  H_though[H]-->T_thought[T]-->end_thought[END]

  H-->R_through[R]-->O_through[O]-->U_through[U]-->G_through[G]-->H_through[H]-->end_through[END]

  U_though[U]-->end_thou[END]
```

Code analysis

- Splits the document into a list of words, and builds each word on the tree character by character, starting at the root each time. For each character, it checks if the correct node exists, creates the node if needed, and then marks the last node of the word as the end of the word.

BFS

Code analysis

- I defined a helper function `get_start_node(prefix)`, that is one of the first things called in both BFS and DFS, and this function just makes sure we're at the correct node to start suggesting from based on the typed prefix. I used the deque data structure to create the FIFO queue that is started from the node given from the helper function. For each child, the prefix is updated and if `child.is_word` becomes true, the new word is added to the suggestions. I made use of the `popleft()` method of the deque to make sure the queue follows the FIFO order.

Your output

- the
- thee
- thou
- that
- thag
- there
- their
- though
- thought
- through

DFS

Code analysis

- This search is almost exactly the same as BFS, as it pulls the start node using the helper function and then begins the search from there. However, instead of using the `popleft()` method, I just use the `pop()` method to create the LIFO order (stack), and I had to reverse the list I iterate through so that it starts from the "top" of the stack.

Your output

- the
- thee
- there
- their
- through
- that
- thag
- thou
- though
- thought

Recursive DFS vs Stack-based DFS

- In recursive, I'd have to create a helper function that recursively called on itself and dove through the tree for me, while in stack-based I could just reuse the queue attribute I created for BFS but pull from the right side instead of the left. I decided to use stack-based as it just seemed more natural to me -

I'm not super comfortable with recursion yet, and I also know Python has like a recursion depth limit, and I didn't want to have to deal with all that. Prof. McNichols also recommended using a stack in OOP, so that gave me another reason.

UCS

Code analysis

- Starts by getting the start node using `get_start_node()`, and initializes priority queue (heapq) with the starting cost at 0. It iterates through the heap and pops the lowest cost node first. If this node completes a word, it's added to the suggestions list. For each child, it updates the prefix and computes the new cost, then pushes this child onto the heap.

Your output

- the
- thou
- thee
- thag
- that
- though
- their
- there
- thought
- through

Experimental

- For BFS, the words come up in order of length, with THE being the shortest word coming first, and THROUGH being the longest coming last. For DFS, you explore each letter first before moving on to the next. For example, you get all the T-H-E... words first (THE, THEE, THERE, THEIR) before you move on to the other letters. For UCS, the output seems a little unordered at first, but looking at the tree you can see that it prioritizes least-cost paths, meaning words with more frequent prefixes are explored first. It explores THE first because it costs the least, and through last because T-H-R... points to nothing else besides THROUGH, making it the most expensive path.